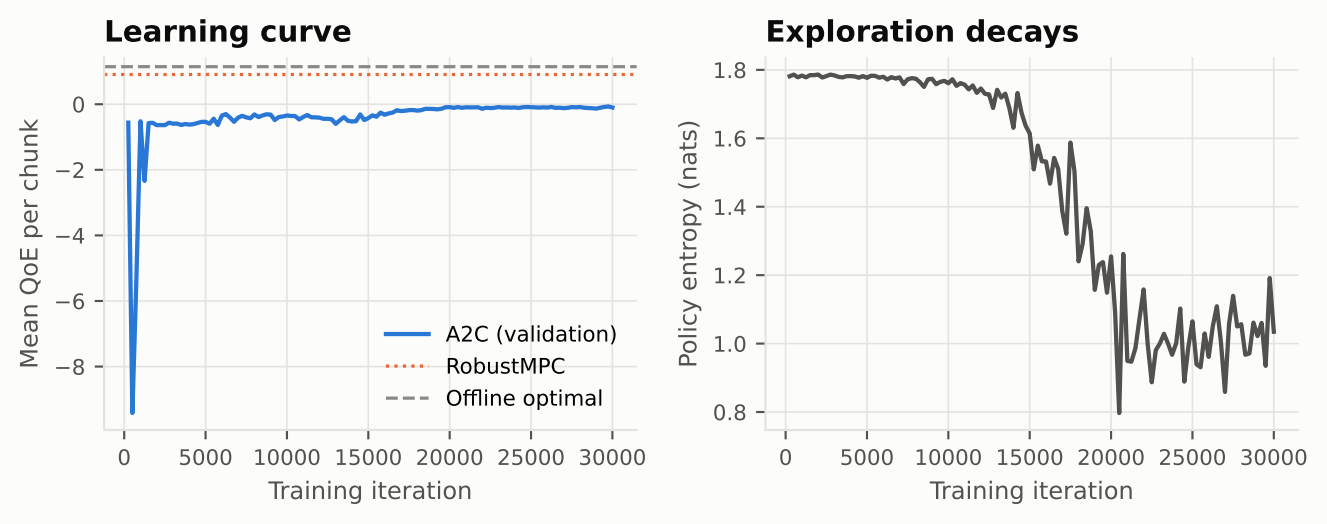
Learning curve
Mean QoE per chunk
0
−2
−4
−6
−8
0
5000
10000
15000
20000
25000
30000
Training iteration
A2C (validation)
RobustMPC
Offline optimal
Exploration decays
Policy entropy (nats)
1.8
1.6
1.4
1.2
1.0
0.8
0
5000
10000
15000
20000
25000
30000
Training iteration