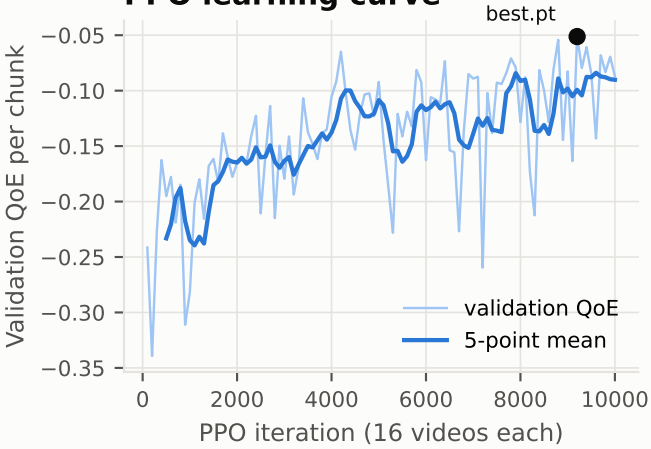
PPO learning curve
best.pt
−0.05
−0.10
−0.15
−0.20
−0.25
−0.30
−0.35
Validation QoE per chunk
validation QoE
5-point mean
0
2000
4000
6000
8000
10000
PPO iteration (16 videos each)

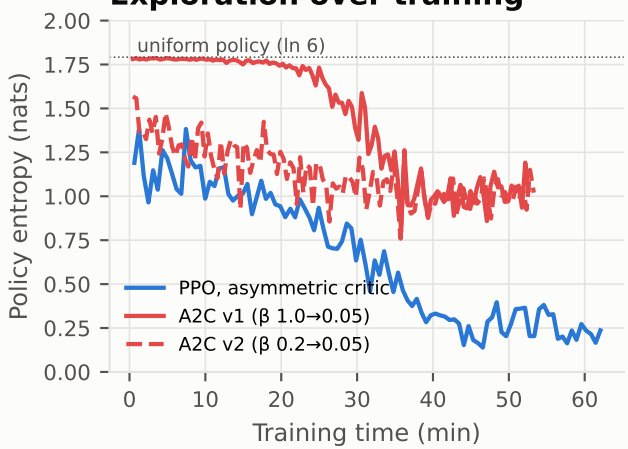
Exploration over training
uniform policy (ln 6)
2.00
1.75
1.50
1.25
1.00
0.75
0.50
0.25
0.00
Policy entropy (nats)
PPO, asymmetric critic
A2C v1 (β 1.0→0.05)
A2C v2 (β 0.2→0.05)
0
10
20
30
40
50
60
Training time (min)